

Lecture 8: Clustering (Part I)

Statistical Methods for Data Science

Yinan Yu

Department of Computer Science and Engineering

Nov 28, 2024

Learning outcome



- Understand the difference between supervised learning and unsupervised learning
- Understand how to apply clustering algorithms to the applications discussed in this lecture
- Be able to implement the K-means algorithm
- Be able to explain the within-cluster sum of squared error (SSE) and the Silhouette score
- Be able to determine K and the best initial guesses using SSE and the Silhouette score

Today

- 1 Introduction
- 2 Mathematical model for clustering
- 3 Preliminary
- 4 Centroid clustering: K-means
- 5 Summary

Clustering

- Input (left): we start with blobs of data points (original data)




- Output (right): we assign each of these data points to a specific group
- Each group is called a **cluster**
- The process of finding these clusters is called **clustering**
- Clustering is widely used for different purposes; clustering algorithm development **does not require expensive annotations**; clustering is **unsupervised**

Application

1. To reduce a large amount of data into fewer data points by, e.g., representing a data set with only a few centroids

Example: you have access to the time people spend on Netflix and the amount of soda they consume everyday; you want to find patterns from this data set




Group these people into clusters and only use the **centroids** for data exploration or as input data for downstream analysis

One important application is the **recommender system**

- Task: find patterns of preferred items from a massive number of users
- Challenge: there are too many users (all data points)
- Solution: we recommend items to users on a cluster level (only the centroids)

Application (cont.)

2. To detect and remove **outliers** - data points significantly distant from any of the clusters (here we assume that the clusters represent the underlying model)



Application (cont.)

3. Image compression



- Each data point is a pixel in the image, i.e. $\mathbf{x} = [red, green, blue] = [x_1, x_2, x_3]$, where $red, green, blue \in [0, 255]$ integers
- Run clustering algorithms in this RGB color space and find K centroids
- Replace each pixel with its closest centroid
- Now we only use $3 \times K$ unique values to represent the image instead of 3×256 values
- In this example, with $K = 10$ centroids, when we save the .png image, we have a reduction from 328.5 kB to 43.4 kB

Today

- 1 Introduction
- 2 Mathematical model for clustering
- 3 Preliminary
- 4 Centroid clustering: K-means
- 5 Summary

Clustering modeling

- Modeling for clustering

$$y = g(x; \theta | h)$$

- Clustering:

- y : **categorical (nominal)**, scalar - each category is called a **cluster**
- x : typically **continuous numerical**; feature vector $\mathbf{x} = [x_1, \dots, x_d]$ (similar to classification problems in lecture 5)
- g : **clustering model**, e.g. K-means, Gaussian mixture models, hierarchical clustering models, etc

There are mainly four categories of clustering models

- **Centroid clustering** (geometry-based)
- **Distribution clustering** (probability-based)
- Density clustering
- Hierarchical clustering
- θ (parameters) and h (hyperparameters) depend on g

Parameter estimation

- Clustering models are **unsupervised learning** algorithms
- In unsupervised learning, the parameters are estimated from an **unlabeled data set**, that is, a data set containing only the feature vectors $\{\mathbf{x}_1, \dots, \mathbf{x}_N\}$, e.g.



where \mathbf{x}_i = pixel values in a picture and the task is to group **similar** ducks into the same cluster

- Clustering tasks do not require annotations - it is cheaper, but also more difficult to evaluate because there are no predefined clusters!
- The **similarity** is not uniquely defined
- In this course, we will look at one commonly used parameter estimation technique called the **Expectation-Maximization (EM)** algorithm

Models in this course

We are going to introduce various categories of clustering techniques; then we focus on two clustering models


- K-means (centroid clustering)
 - **Parameters:** K centroids
 - **Hyperparameters:** the number of centroids K
 - **Parameter estimation:** an iterative method to update the centroids until convergence; this method can be interpreted as a simplified version of the Expectation-Maximization algorithm
- Gaussian mixture models (distribution clustering)
 - **Parameters:** K priors, K Gaussian likelihood (the big two!)
 - **Hyperparameters:** the number of Gaussian components K
 - **Parameter estimation:** the Expectation-Maximization algorithm

Today


- 1 Introduction
- 2 Mathematical model for clustering
- 3 Preliminary**
- 4 Centroid clustering: K-means
- 5 Summary

Let's try something out!

- Generate some data $\{[x_1^1, x_2^1], \dots, [x_1^N, x_2^N]\}$ from a uniform distribution (`np.random.uniform`) 🙄
 - there are **no clusters**



- Run a clustering algorithm



Take one step back: is the data “clusterable”?

- Do you see any clusters in the following plots?




- Figure 1: data is generated from a uniform distribution - no cluster
- Figure 2: data is generated from three different Gaussian distributions - three clusters
- Figure 3: data is generated from two different Gaussian distributions - two clusters
- Figure 4: data is generated from one Gaussian distribution - one cluster
- How to decide if the data is “clusterable”?
 - Need to define what a cluster is
 - Need to define the “null hypothesis”, i.e. the situation where there are no clusters
- There is no ground truth label - there are various ways of defining these prerequisites, which makes it a difficult task!
- Now spend 30 secs staring at the plots and try to think how you can measure if the data set is clusterable

Clustering tendency

The general idea is to **compare** the **data distribution** with a **theoretical distribution with no clustering tendency**!

Let $\mathbf{x}_i = [x_1^i, \dots, x_d^i]$ be a feature vector in this course, when we need to index both the dimension of the feature vector and the data point, we use a superscript to index the data point and a subscript to index the dimension

- For example, we can make a qq-plot to compare the set $\{x_j^1, \dots, x_j^N\}$ and a non-clusterable theoretical probability distribution, e.g. a uniform distribution



- We can repeat this for all dimensions $j = 1, \dots, d$
- But then the question is how to aggregate all these d dimensions? - Not easy!
- **Comparing distributions gets trickier when $d \gg 1$!**
- Approximation and sampling

Distance measure

- Measures how “similar” two data points are
- The most commonly used distance is the Euclidean distance
- Example: let $\mathbf{x} = [x_1, x_2, x_3]$ and $\mathbf{y} = [y_1, y_2, y_3]$ be two feature vectors, the Euclidean distance is defined as

$$d(\mathbf{x}, \mathbf{y}) = \sqrt{(x_1 - y_1)^2 + (x_2 - y_2)^2 + (x_3 - y_3)^2}$$

- A one dimensional value (scalar) computed from high dimensional data (here $d = 3$)
- Pairwise distance
 - Distances between all pairs of data points from two sets
 - Example: let $\{\mathbf{x}_1, \mathbf{x}_2, \mathbf{x}_3\}$ and $\{\mathbf{y}_1, \mathbf{y}_2\}$ be two sets, the pairwise distance is defined as

$$\{d(\mathbf{x}_1, \mathbf{y}_1), d(\mathbf{x}_1, \mathbf{y}_2), d(\mathbf{x}_2, \mathbf{y}_1), d(\mathbf{x}_2, \mathbf{y}_2), d(\mathbf{x}_3, \mathbf{y}_1), d(\mathbf{x}_3, \mathbf{y}_2)\}$$


- Generally useful concept
- Pairwise distance can be used to identify clustering tendency by comparing the **distribution of pairwise distances computed from data** to pairwise distances computed from a distribution without clustering tendency, e.g. a **uniform distribution**

Today

- 1 Introduction
- 2 Mathematical model for clustering
- 3 Preliminary
- 4 Centroid clustering: K-means**
- 5 Summary

K-means

- **Data:** d dimensional feature vector \mathbf{x} (in this example $d = 2$)



- **Target** (the coloring in this image - for each \mathbf{x} , we would like to assign a color to it):

$$y = \arg \min_{k \in \{1, \dots, K\}} \text{dist}(\mathbf{x}, \boldsymbol{\mu}_k)$$

where $\text{dist}(\cdot, \cdot)$ is a distance measure; in this course, we use the Euclidean distance (cf. page 16)

- **Parameters:** K centroids $\hat{\boldsymbol{\mu}}_k$
- **Hyperparameters:** K
- **Parameter estimation:** an iterative method to update the centroids until convergence
- It is a **hard clustering** technique - one data point \mathbf{x} is assigned to only one cluster y

K-means parameter estimation algorithm to find μ_k

- Initialization: **Randomly choose K centroids μ_k** for $k = 1, \dots, K$, e.g. randomly choose K data points from the data set \mathcal{X}
- Repeat the two steps below until **convergence**, e.g. $\hat{\mu}_k$ does not change anymore
 - Step 1: For all $i = 1, \dots, N$, assign \mathbf{x}_i to a cluster \hat{k}_i by computing

$$\hat{k}_i = \arg \min_{k \in \{1, \dots, K\}} \text{dist}(\mathbf{x}_i, \hat{\mu}_k)$$

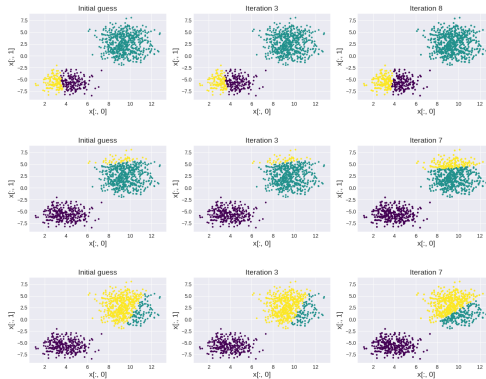
- Step 2: Let \mathcal{X}_k be the set of all \mathbf{x}_i assigned to cluster k and N_k is the size of \mathcal{X}_k , compute

$$\hat{\mu}_k \leftarrow \frac{1}{N_k} \sum_{\mathbf{x}_j \in \mathcal{X}_k} \mathbf{x}_j$$

- There is some **randomness** in the algorithm - we should always be careful when there is randomness

K-means initial guess

Different initializations result in different clusters



A typical solution is to run the algorithm multiple times with different initial points and aggregate the results

K-means parameter estimation pseudocode

- 1: Given a data set $\mathcal{X} = \{\mathbf{x}_1, \dots, \mathbf{x}_N\}$
- 2: Randomly choose K data points from \mathcal{X} as the centroids μ_k for $k = 1, \dots, K$
- 3: **while** true **do**
- 4: Assign \mathbf{x}_i to the closest μ_k for all $i = 1, \dots, N$
- 5: For all $k = 1, \dots, K$, compute μ_k^{new} as the center of all \mathbf{x}_i assigned to cluster k
- 6: **if** $\mu_k^{new} == \mu_k$ for all k **then**
- 7: **break**
- 8: **else**
- 9: $\mu_k \leftarrow \mu_k^{new}$
- 10: **end if**
- 11: **end while**

K-means: pros and cons

- Pros:
 - Convergence guaranteed
 - Easy to implement
 - Scale to large data sets
- Cons - **potential improvement**:
 - Need to choose the hyperparameter K manually - **gradually increase K and monitor the loss during parameter estimation**
 - Dependence on random initial values - **multiple initial values**
 - Do not work well on very high dimensional data - **apply dimensionality reduction techniques before clustering**
 - Not robust to outliers - **try to remove “obvious” outliers before clustering**

Two main challenges for K-means

- **Challenges:**

- How to choose the hyperparameter K ?
- K-means is sensitive to the initialization of $\hat{\mu}_k$ for $k = 1, \dots, K$

- **Solution** (for both challenges):

- Choose a set of **candidate values**, e.g. for the first problem, we can choose $K \in \{1, \dots, 10\}$; for the second problem, we can randomly select 100 different initial guesses for $\hat{\mu}_k$
- For each of these **candidate values**, we run the K-means algorithm to estimate the parameters and evaluate the **quality** of the clusters produced by these parameters
- Choose the **candidate value** that gives the best **quality**

- **Quality** evaluation criteria

- Within-cluster sum of squared errors (SSE)
- Silhouette score

Cluster quality evaluation criterion 1: SSE

1. **Within-cluster sum of squared errors (SSE)**: defined as the summation of the distances from all the data points to their closest centroid

$$SSE = \sum_{k=1}^K \sum_{\mathbf{x} \in C_k} \text{dist}(\mathbf{x}, \hat{\boldsymbol{\mu}}_k)^2 \quad (1)$$

where C_k denote cluster k ; $\text{dist}(\cdot, \cdot)$ is a distance measure (**Euclidean distance**: $\text{dist}(\mathbf{x}, \hat{\boldsymbol{\mu}}_k)^2 = (\mathbf{x} - \hat{\boldsymbol{\mu}}_k)^T (\mathbf{x} - \hat{\boldsymbol{\mu}}_k)$ for column vectors \mathbf{x} and $\hat{\boldsymbol{\mu}}_k$)

Example:

- Given $\mathbf{x}_1 = [x_1^1, x_2^1]$, $\mathbf{x}_2 = [x_1^2, x_2^2]$, $\mathbf{x}_3 = [x_1^3, x_2^3]$, $\mathbf{x}_4 = [x_1^4, x_2^4]$, where $\mathbf{x}_1, \mathbf{x}_2 \in$ cluster 1 with centroid $\boldsymbol{\mu}_1 = [\mu_1^1, \mu_2^1]$; $\mathbf{x}_3, \mathbf{x}_4 \in$ cluster 2 with centroid $\boldsymbol{\mu}_2 = [\mu_1^2, \mu_2^2]$
- The SSE is computed as

$$\begin{aligned} SSE &= \text{distance in cluster 1} + \text{distance in cluster 2} \\ &= \underbrace{(x_1^1 - \mu_1^1)^2 + (x_2^1 - \mu_2^1)^2}_{\text{dist}(\mathbf{x}_1, \boldsymbol{\mu}_1)^2} + \underbrace{(x_1^2 - \mu_1^1)^2 + (x_2^2 - \mu_2^1)^2}_{\text{dist}(\mathbf{x}_2, \boldsymbol{\mu}_1)^2} \\ &\quad + \underbrace{(x_1^3 - \mu_1^2)^2 + (x_2^3 - \mu_2^2)^2}_{\text{dist}(\mathbf{x}_3, \boldsymbol{\mu}_2)^2} + \underbrace{(x_1^4 - \mu_1^2)^2 + (x_2^4 - \mu_2^2)^2}_{\text{dist}(\mathbf{x}_4, \boldsymbol{\mu}_2)^2} \end{aligned}$$

Cluster quality evaluation criterion 1: SSE (cont.)

1. Within-cluster sum of squared errors (SSE) (cont.):

$$SSE = \sum_{k=1}^K \sum_{\mathbf{x} \in C_k} \text{dist}(\mathbf{x}, \hat{\mu}_k)^2$$

- SSE is essentially an error term: we want SSE to be small - choose the K value that minimizes SSE ?
- No can do! $SSE \rightarrow 0$ for $K \rightarrow N$, i.e. when every data point is their own centroid, $SSE = 0$, which is not optimal - we can't simply choose the K value that corresponds to the smallest SSE
- Instead, the best K is defined as the **elbow** point of the SSE (instead of the minimum), i.e. **the point with the maximum curvature** - find the largest K where the SSE does not go down significantly by further increasing K
- This method is also called the **elbow method** (in Python, you can find a library to compute the elbow point)

Cluster quality evaluation criterion 2: Silhouette score

2. **Silhouette score S** : the idea is that a good clustering should have **compact clusters** with a **large separation between different clusters**. This is characterized by the **within-cluster distance** and **between-cluster distance**

Cluster quality evaluation criterion 2: Silhouette score (cont.)

2. Silhouette score S (cont.):

Example: given data $x_1, x_2, x_3 \in C_1$, $x_4, x_5 \in C_2$, $x_6, x_7 \in C_3$; $K = 3$; C_k denotes the set of cluster k ; $|C_k|$ is the cardinality (size) of the set C_k

- **Within-cluster distance:** measures how data points scatter in relation to x_i within its own cluster; let x_i be a data point from cluster k ,

$$a_i = \frac{1}{|C_k| - 1} \sum_{x_j \in C_k \text{ and } j \neq i} \text{dist}(x_i, x_j)$$

In this example, let $i = 1$, $x_1 \in C_1$; there are $|C_1| = 3$ data points in cluster 1

$$a_1 = \frac{1}{3 - 1} (\text{dist}(x_1, x_2) + \text{dist}(x_1, x_3))$$

- **Between-cluster distance:** measures how data points scatter in relation to x_i when these data points are from other clusters

$$b_i = \min_{k' \neq k, k' \in \{1, \dots, K\}} \frac{1}{|C_{k'}|} \sum_{x_j \in C_{k'}} \text{dist}(x_i, x_j)$$

In the example, $|C_2| = |C_3| = 2$

$$b_1 = \min \left(\frac{1}{2} (\text{dist}(x_1, x_4) + \text{dist}(x_1, x_5)), \frac{1}{2} (\text{dist}(x_1, x_6) + \text{dist}(x_1, x_7)) \right)$$

Cluster quality evaluation criterion 2: Silhouette score (cont.)

2. Silhouette score S (cont.):

- Silhouette score for one data point \mathbf{x}_i :

$$S_i = \begin{cases} \frac{b_i - a_i}{\max(a_i, b_i)}, & \text{if } |C_k| > 1 \\ 0, & \text{if } |C_k| = 1 \end{cases}$$

A large S_i indicates a compact cluster k in relation to \mathbf{x}_i and a large distance from \mathbf{x}_i to clusters other than k

- Silhouette score for the data set


$$S = \frac{1}{N} \sum_{i=1}^N S_i, \quad S \in [-1, 1]$$

- A large Silhouette score indicates a good clustering quality

Example - choose K


Clusters with equal variance ($K = 8$)

- SSE: $K = 8$
- Silhouette score: $K = 8$



Overlapping clusters with unequal variances ($K = 8$)

- SSE: $K = 6$
- Silhouette score: $K = 8$; but $K = 6$ and $K = 8$ have similar Silhouette scores

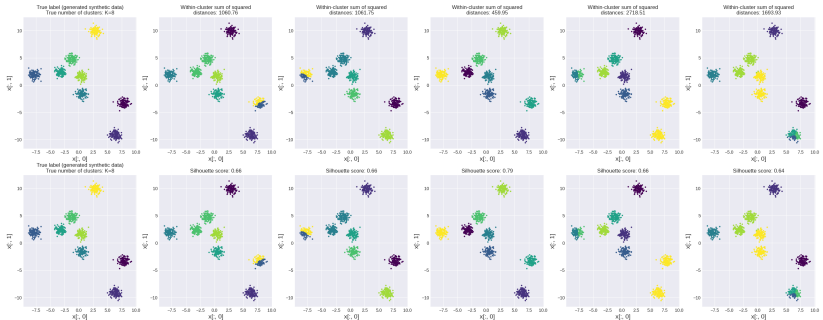


Example - choose initial guess

- Each column corresponds to a different initialization
- For a given K , choose the initialization that gives the smallest SSE or largest Silhouette score

Clusters with equal variance ($K = 8$)

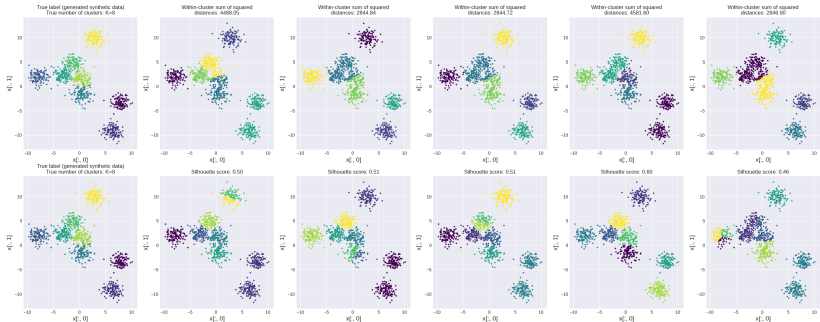
- SSE: $K = 8$
- Silhouette score: $K = 8$



Example - choose initial guess (cont.)

Overlapping clusters with unequal variances ($K = 8$)

- SSE: $K = 6$
- Silhouette score: $K = 8$



Today

- 1 Introduction
- 2 Mathematical model for clustering
- 3 Preliminary
- 4 Centroid clustering: K-means
- 5 Summary**

Summary

So far:


- Data types and data containers
- Descriptive data analysis: descriptive statistics, visualization
- Probability distributions, events, random variables, PMF, PDF, parameters
- CDF, Q-Q plot, how to compare two distributions (data vs theoretical, data vs data)
- Modeling
- Parameter estimation: maximum likelihood estimation (MLE) and maximum a posteriori estimation (MAP)
- Classification, multinomial naive Bayes classifier, Gaussian naive Bayes classifier
- Clustering, clustering tendency
- Centroid clustering, k-means, parameter estimation, SSE, Silhouette score

Next:

- Gaussian Mixture Models (GMMs)

Before next lecture:

- Gaussian distribution
- The Bayes' rule



You will never get rid of me!